

# Revisiting Link Prioritization for Efficient Traversal in Structured Decentralized Environments

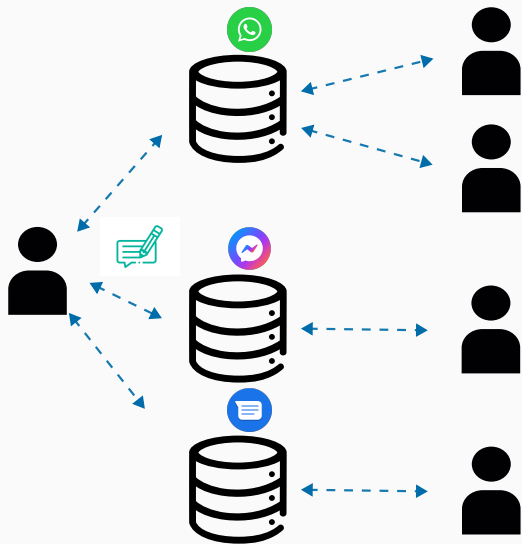
---

Ruben Eschauzier, Ruben Taelman, Ruben Verborgh

October 28, 2025

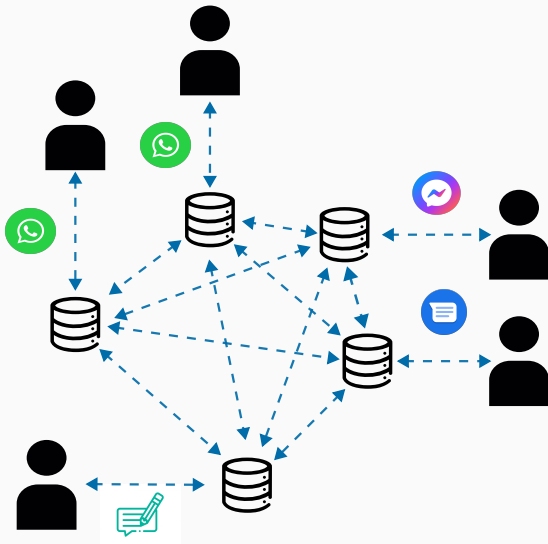
Department of Electronics and Information Systems, Ghent University Imec

# The Need for Decentralized Personal Data Storage



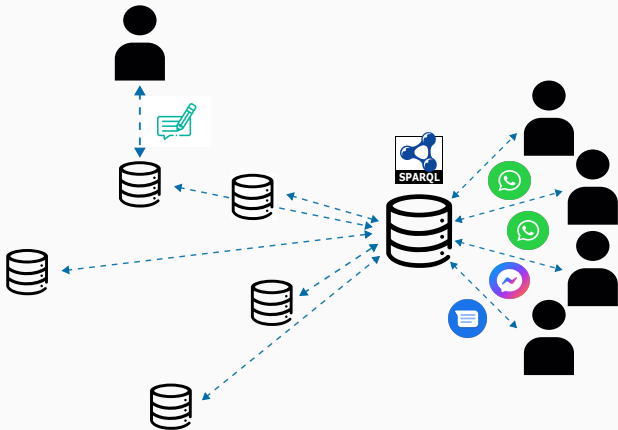
- Each application has its own data
- Stifles innovation
- Causes vendor lock-in

# The Need for Decentralized Personal Data Storage



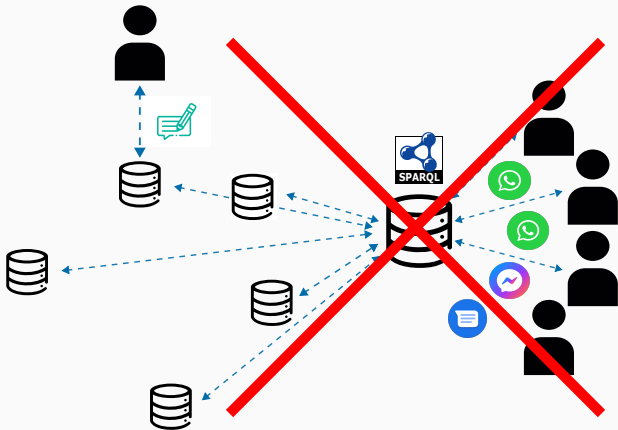
- Each application uses common data storage
- Easy to switch vendors
- Promotes innovation

# The Problem with Centrally Aggregating and Querying




- Why not aggregate data and query it?
- Difficult in case of personal data due to privacy concerns

# The Problem with Centrally Aggregating and Querying




- Why not aggregate data and query it?
- Difficult in case of personal data due to privacy concerns

## Link Traversal: Seed Document




- Link Traversal starts from seed documents (URIs)
- These are provided by the user or in the query.

## Link Traversal: Traversal



- New URIs are extracted from the seed document
- URIs are extracted in accordance with reachability criteria


## Link Traversal: Traversal



- New URIs are dereferenced and the process is repeated



## Link Traversal: Termination




- This continues until all links are dereferenced
- Continuously produces results
- Can enforce fine-grained (document-level) access-control

# Query Optimization for Link Traversal

```
SELECT * WHERE {  
  <seedUri> <ex:p1> ?o1.  
  <seedUri> <ex:p2> ?o2.  
  ...  
}
```

Extract  
Seed document




The query (partly) determines:


- The queried data
- The topology of the queried data
- The query-relevant documents

Result: limited prior knowledge for query optimization

# Query Optimization for Link Traversal: Traditional Query Planning

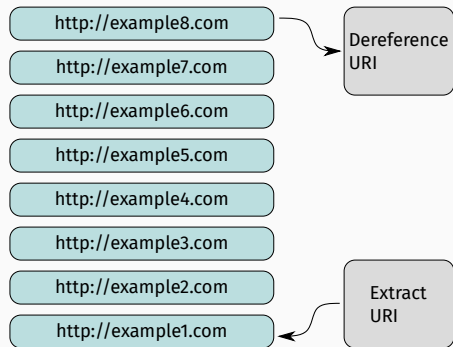


VS




- Query optimization for link traversal involves traditional (zero knowledge) query planning
- Assume an optimize-then-execute

## Query Optimization for Link Traversal: Traversal optimization



- URIs are put into a link queue
- Link traversal uses a FiFo queue by default

## Query Optimization for Link Traversal: Traversal optimization



### Optimizations:

- Prioritize query-relevant URIs
- Prune irrelevant URIs
- Pruning can lead to missing results without prior knowledge


# Problem Statement

Investigate the performance of link prioritization algorithms in literature in new structured decentralized environments


# Problem Statement

- These algorithms are implemented in an engine that is no longer maintained
- The baseline performance depends on design choices orthogonal to the prioritization algorithm
- The baseline should measure marginal prioritization performance

## $R^3$ Metric: Introduction



**Figure 1:** Actual traversal



**Figure 2:** Optimal traversal: minimal cost path

$$R^3 = \frac{5}{13} = 0.38$$



### Definition

$$R^3 = \frac{C(X)}{C(O_{\mathcal{T}})}$$

- $X$ : the **optimal traversal order** (minimal-cost path)
- $O_{\mathcal{T}}$ : traversal order produced by the link prioritization algorithm
- $C(\cdot)$ : The total cost of all arcs in the traversal path

### Interpretation

- $R^3 = 1$  perfect match with the optimal traversal
- $R^3 < 1$  deviation from optimal performance
- **Higher values indicate better prioritization quality.**

- Directed graph steiner tree over traversed directed graph  $G = (V, A)$  with root  $r$
- Query-relevant documents  $D_T$  serve as terminals  $T \subseteq V$
- Find minimum cost sub-graph  $X = (V', A')$  starting at root  $r$  and spanning all vertexes  $T$ .
- With cost:  $C(X) = \sum_{a \in A'} c(a)$ , with  $c(a)$  the cost of an edge in the topology

# Investigated Structured Decentralized Environment

## **Simulated Solid Environment: SolidBench (Taelman and Verborgh 2023)**

- Data is stored in vaults
- Vaults use document-centric structure
- Vaults simulate a social media application

### **Data attributes**

- Generates 158,233 RDF files
- Across 1,531 data vaults
- containing a total of 3,556,159 triples.

### **Analysed queries**

- Use discover queries


## Non-adaptive

- Breadth-first (default), depth-first, random prioritization

## Graph-based

- In-degree, PageRank score


## Prioritization algorithms: Result Contribution-Based (RCC)



- Nodes adaptively scored according to the number of results they contribute to
- Called *rcc-1*, *rcc-2*, *rel-1*, *rel-2*.

## Prioritization algorithms: Intermediate Results-based

```
SELECT ?person ?project ?topic
WHERE {
  ?person ex:worksFor ?company .
  ?person ex:worksOn ?project .
  ?project ex:hasTopic ?topic .
}
```



- Nodes adaptively scored according to their contribution to intermediate results
- *IS* with initial priorities of 0, while *ISdcr* sets priority to the priority of the parent node - 1

# Prioritization algorithms

## Hybrid

- Multiply intermediate result and RCC-based scoring functions
- *is-rcc1*, *is-rcc2*, *is-rel1*, *is-rel2*


## TypeIndex

- TypeIndex points to location for resource of specific type
- Prioritize TypeIndex

## Oracle


- Compute RCC in hindsight
- Scores are propagated through the shortest path
- Serves as optimal performance oracle

# Prioritization Has Limited Impact on Result Arrival Time







# Prioritization Has Limited Impact on Result Arrival Time



# Prioritization Has Limited Impact on Result Arrival Time



# Prioritization Has Limited Impact on Result Arrival Time



## Prioritization Has Limited Impact on Result Arrival Time





	1st		Cmpl		R <sup>3</sup>	
	better	worse	better	worse	better	worse
depth-first	<u>15.7</u>	<u>17.1</u>	<u>14.3</u>	17.1	7.5	12.5
random	4.3	68.6	5.7	74.3	<u>15.0</u>	15.0
is-rcc-1	4.3	57.1	5.7	60.0	7.5	<u>5.0</u>
type-index	<u>15.7</u>	18.6	10.0	<u>15.7</u>	12.5	20.0
oracle	<u><b>18.6</b></u>	<u><b>11.4</b></u>	<u><b>18.6</b></u>	<u><b>12.9</b></u>	<u><b>25.0</b></u>	<u><b>2.5</b></u>

**Table 1:** Percentage of queries for which an algorithm performs at least 10% better or worse than the breadth-first baseline.

- No algorithm improves performance, and the oracle only marginally improves performance

## Link Prioritization in Structured Decentralized Environments

- Link prioritization in Solid environment will not significantly improve query performance
- Research should instead focus on query planning Hanski et al. (2025) or traversal pruning Tam et al. (2024)

-  Hanski, J. et al. (2025). **“Link Traversal over Decentralised Environments using Restart-Based Query Planning”**. In: *International Conference on Web Engineering*.
-  Hartig, O. and M. T. Özsu (2016). **“Walking without a map: Ranking-based traversal for querying linked data”**. In: *International Semantic Web Conference*. Springer, pp. 305–324.
-  Taelman, R. and R. Verborgh (2023). **“Link traversal query processing over decentralized environments with structural assumptions”**. In: *International Semantic Web Conference*. Springer, pp. 3–22.
-  Tam, B.-E. et al. (2024). **“Opportunities for Shape-based Optimization of Link Traversal Queries”**. In: *arXiv preprint arXiv:2407.00998*.